

# DeBERTa: Decoding-enhanced BERT with Disentangled Attention

Radek Bartyzal

GLAMI AI

2. 3. 2021

# Motivation

Paper is by Microsoft Research from 2020/2021.

- improved BERT
- models available through Hugging Face
- new SOTA in Language Understanding and Lang. Generation downstream tasks as of January 6, 2021
- faster pre-training = less steps, but less time?

# BERT

Transformer +

# DeBERTa changes to BERT

**disentangled attention mechanism:** treat content and position information separately

**relative positional encodings** passed to each layer

**enhanced mask decoder:** provide absolute positional encodings right before the last softmax layer decoding the masked words

**Scale invariant Fine Tuning:** perturb normalized word embeddings = adversarial training = better generalization

# Classic attention mechanism

$$Q = HW_q, K = HW_k, V = HW_v, A = \frac{QK^\top}{\sqrt{d}}$$
$$H_o = \text{softmax}(A)V$$

- before the first layer word vectors are summed with positional vectors
- these combined = entangled content vectors are then passed through Transformer layers
- each layer's attention creates Key, Query and Value vectors from these combined vectors

# Disentangled attention mechanism

disentangled attention = at each layer, create separate Key and Query vectors =  $K^r$  and  $Q^r$  from relative positional vectors

$$Q_c = HW_{q,c}, K_c = HW_{k,c}, V_c = HW_{v,c}, Q_r = PW_{q,r}, K_r = PW_{k,r}$$

$$\tilde{A}_{i,j} = \underbrace{Q_i^c K_j^{c\top}}_{\text{(a) content-to-content}} + \underbrace{Q_i^c K_{\delta(i,j)}^r\top}_{\text{(b) content-to-position}} + \underbrace{K_j^c Q_{\delta(j,i)}^r\top}_{\text{(c) position-to-content}}$$

$$H_o = \text{softmax}\left(\frac{\tilde{A}}{\sqrt{3d}}\right)V_c$$

- there is still only one Value vector created from content vector
- content vectors are the ones transformed by each layer, positional vectors stay the same
- final attention matrix is sum of attention matrices of possible combinations of Pos. and Content Q/Ks

# Disentangled attention mechanism

$P_{i,j}$  = encoding of a position of  $i$  relative to  $j$

$$\begin{aligned} A_{i,j} &= \{H_i, P_{i|j}\} \times \{H_j, P_{j|i}\}^\top \\ &= H_i H_j^\top + H_i P_{j|i}^\top + P_{i|j} H_j^\top + P_{i|j} P_{j|i}^\top \end{aligned}$$

Possible attention matrices = combinations of P, C:

- C to C = classic attention
- C to P = what relative positions are important to this word
- P to C = what words are important to these relative positions
- P to P = what relative position are important to this relative position  
= does not make sense = not used

# Disentangled attention: Efficient implementation

- an input sequence of length  $N$
- requires a space complexity of  $O(N^2d)$  to store the relative position embedding for each token
- set the maximum relative distance  $k$  to 512 for pre-training
- embeddings of all possible relative positions are always a subset of  $K_r \in R^{2k \times d} \implies$  we can reuse  $K_r$  in the attention calculation for all the queries
- we do not need to allocate memory to store a relative position embedding for each query and thus reduce the space complexity to  $O(kd)$  (probably  $O(2kd)$ ) (for storing  $K_r$  and  $Q_r$ )



# Relative positional encodings

Denote  $k$  as the maximum relative distance,  $\delta(i, j) \in [0, 2k)$  as the relative distance from token  $i$  to token  $j$ , which is defined as:

$$\delta(i, j) = \begin{cases} 0 & \text{for } i - j \leq -k \\ 2k - 1 & \text{for } i - j \geq k \\ i - j + k & \text{others.} \end{cases} \quad (3)$$

- the positional vectors are relative = [..., -2, -1, 0, 1, 2, ...]
- the positional encodings are shared between layers
- max relative distance =  $k = 512$ , after that just pad: [-2, -2, -1, 0]

# Scale invariant Fine Tuning = SiFT


- The model is regularized: model produces the same output distribution as it produces on an adversarial perturbation of that example.
- perturbation is applied to the word embedding instead of the word sequence
- However, the value ranges (norms) of the embedding vectors vary among different words and models.
- variance gets larger for bigger models with billions of parameters, leading to some instability of adversarial training.
- SiFT improves the training stability by applying the perturbations to the normalized word embeddings.
- Specifically, when fine-tuning DeBERTa to a downstream NLP task in our experiments, SiFT first normalizes the word embedding vectors into stochastic vectors, and then applies the perturbation to the normalized embedding vectors.
- We find that the normalization substantially improves the performance of the fine-tuned models.

# Ablation study: everything helps


Model	MNLI-m/mm Acc	SQuAD v1.1 F1/EM	SQuAD v2.0 F1/EM	RACE Acc
BERT <sub>base</sub> Devlin et al. (2019)	84.3/84.7	88.5/81.0	76.3/73.7	65.0
RoBERTa <sub>base</sub> Liu et al. (2019c)	84.7/-	90.6/-	79.7/-	65.6
XLNet <sub>base</sub> Yang et al. (2019)	85.8/85.4	-/-	81.3/78.5	66.7
RoBERTa-ReImp <sub>base</sub>	84.9/85.1	91.1/84.8	79.5/76.0	66.8
DeBERTa <sub>base</sub>	<b>86.3/86.2</b>	<b>92.1/86.1</b>	<b>82.5/79.3</b>	<b>71.7</b>
-EMD	86.1/86.1	91.8/85.8	81.3/78.0	70.3
-C2P	85.9/85.7	91.6/85.8	81.3/78.3	69.3
-P2C	86.0/85.8	91.7/85.7	80.8/77.6	69.6
-(EMD+C2P)	85.8/85.9	91.5/85.3	80.3/77.2	68.1
-(EMD+P2C)	85.8/85.8	91.3/85.1	80.2/77.1	68.5

Table 5: Ablation study of the DeBERTa base model.

# Pre-training needs less steps than RoBERTa



(a) Results on MNLI development



(b) Results on SQuAD v2.0 development

Figure 1: Pre-training performance curve between DeBERTa and its counterparts on the MNLI and SQuAD v2.0 development set.

# Results: GLUE

Model	CoLA Mcc	QQP F1/Acc	MNLI-m/mm Acc	SST-2 Acc	STS-B Corr	QNLI Acc	RTE Acc	MRPC Acc	Avg.
BERT <sub>large</sub>	60.6	91.3	86.6/-	93.2	90.0	92.3	70.4	88.0	84.05
RoBERTa <sub>large</sub>	68.0	92.2	90.2/90.2	96.4	92.4	93.9	86.6	90.9	88.82
XLNet <sub>large</sub>	69.0	92.3	90.8/90.8	<b>97.0</b>	92.5	94.9	85.9	90.8	89.15
ELECTRA <sub>large</sub>	69.1	<b>92.4</b>	90.9/-	96.9	92.6	95.0	88.0	90.8	89.46
DeBERTa <sub>large</sub>	70.5	92.3	<b>91.1/91.1</b>	96.8	92.8	<b>95.3</b>	88.3	<b>91.9</b>	<b>90.00</b>

Table 1: Comparison results on the GLUE development set.

We use 6 DGX-2 machines (96 V100 GPUs) to train the models. A single model trained with 2K batch size and 1M steps takes about 20 days.

# Results: Comparison with similar sized models

Model	MNLI-m/mm Acc	SQuAD v1.1 F1/EM	SQuAD v2.0 F1/EM	RACE Acc	ReCoRD F1/EM	SWAG Acc	NER F1
BERT <sub>large</sub>	86.6/-	90.9/84.1	81.8/79.0	72.0	-	86.6	92.8
ALBERT <sub>large</sub>	86.5/-	91.8/85.2	84.9/81.8	75.2	-	-	-
RoBERTa <sub>large</sub>	90.2/90.2	94.6/88.9	89.4/86.5	83.2	90.6/90.0	89.9	93.4
XLNet <sub>large</sub>	90.8/90.8	95.1/89.7	90.6/87.9	85.4	-	-	-
Megatron <sub>336M</sub>	89.7/90.0	94.2/88.0	88.1/84.8	83.0	-	-	-
DeBERTa <sub>large</sub>	<b>91.1/91.1</b>	<b>95.5/90.1</b>	<b>90.7/88.0</b>	<b>86.8</b>	<b>91.4/91.0</b>	<b>90.8</b>	<b>93.8</b>
ALBERT <sub>xxlarge</sub>	90.8/-	94.8/89.3	90.2/87.4	86.5	-	-	-
Megatron <sub>1.3B</sub>	90.9/91.0	94.9/89.1	90.2/87.1	87.3	-	-	-
Megatron <sub>3.9B</sub>	91.4/91.4	95.5/90.0	91.2/88.5	89.5	-	-	-

Table 2: Results on MNLI in/out-domain, SQuAD v1.1, SQuAD v2.0, RACE, ReCoRD, SWAG, CoNLL 2003 NER development set. Note that missing results in literature are signified by “-”.

## Results: Scale up to 1.5B parameters

Model	BoolQ Acc	CB F1/Acc	COPA Acc	MultiRC F1a/EM	ReCoRD F1/EM	RTE Acc	WiC Acc	WSC Acc	Average Score
RoBERTa <sub>large</sub>	87.1	90.5/95.2	90.6	84.4/52.5	90.6/90.0	88.2	69.9	89.0	84.6
NEXHA-Plus	87.8	94.4/96.0	93.6	84.6/55.1	90.1/89.6	89.1	74.6	93.2	86.7
T5 <sub>11B</sub>	91.2	93.9/96.8	94.8	88.1/63.3	94.1/93.4	92.5	76.9	93.8	89.3
T5 <sub>11B</sub> +Meena	<b>91.3</b>	<b>95.8/97.6</b>	97.4	88.3/63.0	94.2/93.5	92.7	<b>77.9</b>	95.9	90.2
Human	89.0	95.8/98.9	100.0	81.8/51.9	91.7/91.3	93.6	80.0	100.0	89.8
DeBERTa <sub>1.5B</sub>	90.4	94.9/97.2	96.8	<b>88.2/63.7</b>	<b>94.5/94.1</b>	<b>93.2</b>	76.4	<b>95.9</b>	89.9
DeBERTa <sub>Ensemble</sub>	90.4	95.7/97.6	<b>98.4</b>	88.2/63.7	94.5/94.1	93.2	77.5	95.9	<b>90.3</b>

Table 6: SuperGLUE test set results scored using the SuperGLUE evaluation server. All the results are obtained from <https://super.gluebenchmark.com> on January 6, 2021.

- we share the projection matrices of relative position embedding  $W_{k,r}$ ,  $W_{q,r}$  with  $W_{k,c}$ ,  $W_{q,c}$ , respectively, in all attention layers = not treating the content/positional information differently?

# Conclusion

The single 1.5B-parameter DeBERTa model substantially outperforms T5 with 11 billion parameters on the SuperGLUE benchmark by 0.6%(89.3% vs. 89.9%)



# Sources

1. He, Pengcheng, et al. "Deberta: Decoding-enhanced bert with disentangled attention." arXiv preprint arXiv:2006.03654 (2020).  
<https://arxiv.org/abs/2006.03654>