

Perceiver IO: A General Architecture for Structured Inputs & Outputs

by DeepMind

GLAMI AI
Radek Bartyzal
7. 9. 2021

Motivation

Current state:


- NN architectures usually differ based on the input modality = audio, image, video, ...
- e.g. CNNs and Visual Transformers rely on image specific locality assumption
 - => convolutions, splitting the image into a grid
- Perceiver can only predicts logits

Author's goal:

- create competitive architecture without relying on these assumptions
 - => attend to the individual pixels in an image
 - => use same architecture for audio, video, 3D point cloud
- enable larger inputs to Transformers
- enable any output


Regular Transformers

- transform a set of tokens into another set of tokens of same length = 1 transformer layer
- core of the transformer layer = self-attention



Self Attention

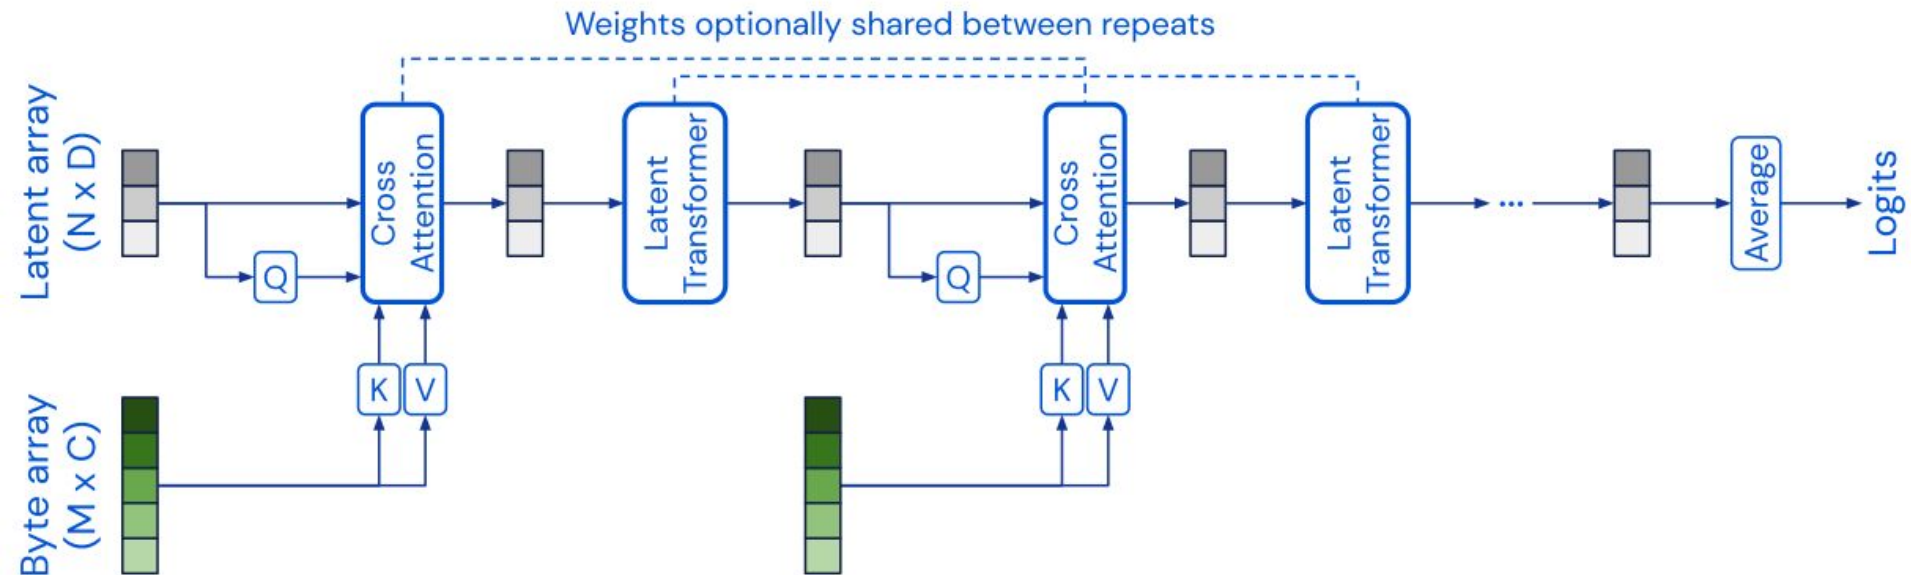
- M input tokens = e.g. word embeddings
- from each input token, create Query, Key, Value vectors
- dot product Query * Key vectors => attention matrix of MxM
- dot product Att matrix * Value vectors => new set of token vectors
- => $O(M^2)$ space, $O(M^2 * d)$ time



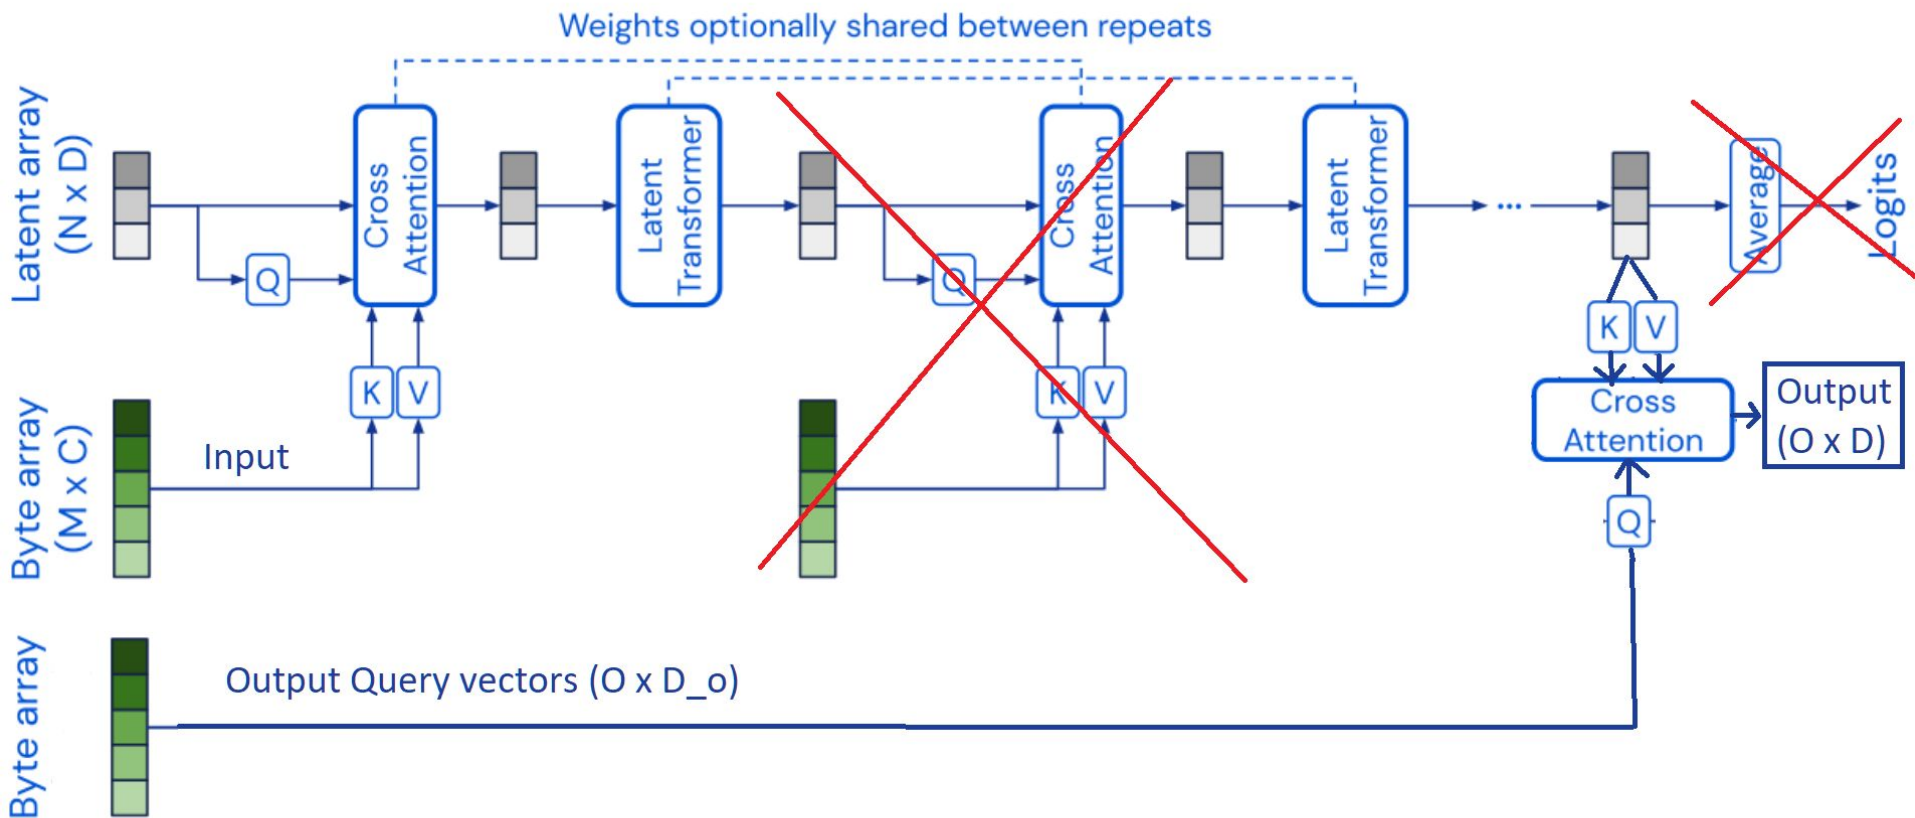
Perceiver architecture

- 224x224 image = 50 000 pixels = M
- \Rightarrow impossible to calculate self-attention with $O(50000*50000)$
- \Rightarrow use **cross-attention** to inject the input information into the latent representation
- \Rightarrow attention from N latent vectors to M input tokens = $O(N*M)$, $N \ll M$
- follow the cross-attention by regular self-attention on the latent space = $O(N*N)$
- repeat blocks of [cross-attention->regular-latent-attention]

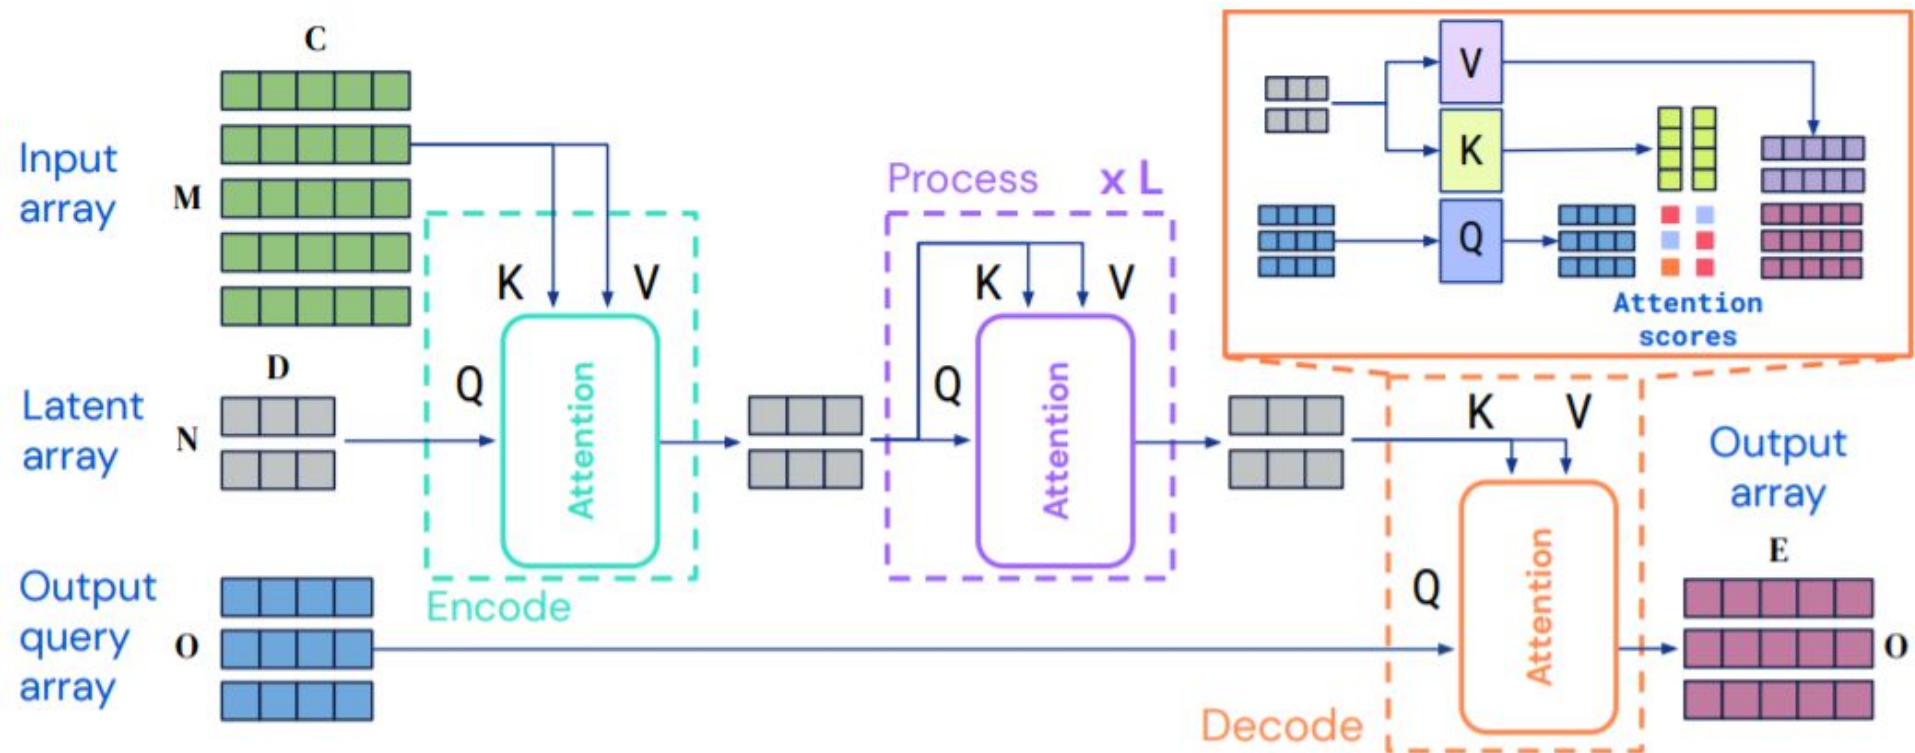
Perceiver architecture



Perceiver IO



Perceiver IO



Perceiver IO


- autoencoder like structure
 - encode = cross-attention of: input -> latent
 - decode = cross-attention of: latent -> output
 - don't repeatedly pass in input = does not help that much
 - => not a RNN anymore :(
-
- positional encoding = domain specific = hand crafted
 - 1D = sequence = words
 - 2D = image
 - 3D = 3D point cloud
 - 4D = video

Output Query vectors = decoder vectors


- Can be:
 - learned
 - hand crafted
 - simple function of the input
- number of decoder vectors = number of outputs = dynamic
- output dimension is fixed = task specific = e.g. 1000 for imagenet
- decode query vectors = task specific = hand crafted **structure**
 - simple classification => single learned embedding = token = query vector
 - word prediction => position in sentence + <predict word task token>
 - optical flow of pixel => x,y coords of pixel + <optical flow task token>
 - Starcraft unit information => embedding of unit description
 - task specification token is needed only for multi-task models

Decoder vectors


Masked language modeling




Classification




StarCraft II




Multi-task classification



Optical flow



Multimodal autoencoding



- task_id = single learned embedding is enough for class. task = imagenet
- position = pos. embedding as query which will return predicted word distribution
- is_label = <learned token of give me label logits = 1000 dim for imagenet>

Tasks

| Modalities | Tasks | Preprocessing | Postprocessing | # Inputs | # Outputs |
|--------------------|-----------------------------|--------------------------|--------------------|---------------------|----------------------|
| Text | Token-level pred. | Tokenization + Embed. | Linear projection | 512×768 | 512×768 |
| Text | Byte-level pred. | Embed. | None | $2,048 \times 768$ | $2,048 \times 768$ |
| Text | Multi-task (8 tasks) | Embed. | None | $2,048 \times 768$ | 8×768 |
| Video | Flow prediction | Conv+maxpool | RAFT upsampling | $22,816 \times 64$ | $11,408 \times 64$ |
| Video+Audio+Label | Auto-encoding | Patch: 1x4x4 Vid, 16 Aud | Linear projections | $50,657 \times 704$ | $803,297 \times 512$ |
| StarCraft Unit Set | Encoding and Classification | Tokenization | Pointer network | 512×256 | 512×128 |
| Image | Classification | None | None | $50,176 \times 3$ | $1 \times 1,000$ |

Table 1: We give details of each of the tasks we use to evaluate Perceiver IO here. The positional and task embeddings appended to inputs for each case are listed in Appendix Table 7.

GLUE

| Model | Tokenization | N (# inputs) | M (# latents) | Depth | Params | FLOPs | Average |
|-----------------------|---------------|----------------|-----------------|-------|--------|-------|-------------|
| BERT Base (test) [21] | SentencePiece | 512 | 512 | 12 | 110M | 109B | 81.0 |
| BERT Base (ours) | SentencePiece | 512 | 512 | 12 | 110M | 109B | 81.1 |
| Perceiver IO Base | SentencePiece | 512 | 256 | 26 | 223M | 119B | 81.2 |
| BERT (matching FLOPs) | UTF-8 bytes | 2048 | 2048 | 6 | 20M | 130B | 71.5 |
| Perceiver IO | UTF-8 bytes | 2048 | 256 | 26 | 201M | 113B | 81.0 |
| Perceiver IO++ | UTF-8 bytes | 2048 | 256 | 40 | 425M | 241B | 81.8 |

Table 2: **Perceiver IO on language:** results on the GLUE benchmark (higher is better). Following [21] we exclude the WNLI task. We use Pearson correlation on STS-B, Matthews correlation on CoLa and accuracy on the remaining tasks.

GLUE

- multitask = 1 model for all tasks
- input concat with special token:
 - 1 for all tasks = shared input token
 - 1 for each = task specific input token
 - = like CLS token = decoder vector
 - a task specific head, a 2-layer MLP, is applied to generate the output logits corresponding to each task
- multitask query = use decoder vectors specific to each task
 - no input token
 - no output 2-layer heads?

| Multitask method | Avg. |
|---------------------------|-------------|
| Single task | 81.0 |
| Shared input token | 81.5 |
| Task specific input token | 81.8 |
| Multitask Query | 81.8 |

Table 3: Multitask Perceiver IO. Results use the same metric as Table 2 (higher is better).

SOTA on Optical flow

- input: 2 images = consecutive frames from video
- output: displacement of each pixel between the 2 images
- PWCNet, RAFT = slow, hand-crafted architectures, very task specific

| Network | Sintel.clean | Sintel.final | KITTI |
|--------------|--------------|--------------|-------------|
| PWCNet [75] | 2.17 | 2.91 | 5.76 |
| RAFT [84] | 1.95 | 2.57 | 4.23 |
| Perceiver IO | 1.81 | 2.42 | 4.98 |

Table 4: Optical Flow evaluated on Sintel [10] and KITTI with average end-point error (EPE) (lower is better). Baselines are reported from [74].

Optical flow

Input:

- concat the frames along channels
- extract 3x3 area around each pixel
- => **3x3 x 3=channels x 2=frames** = 54 values for each pixel
- then concatenate a fixed position encoding to the 54 values

Decoder vector = same embedding as used for input

Audio Video Class reconstruction

- inputs:
 - padded with modality-specific embeddings
 - raw video pixels, audio, one-hot label
- outputs = raw reconstruction of video, audio, one-hot label



Figure 4: Audio-visual autoencoding with 88x compression. Side-by-side: inputs on left, reconstructions right (class label not shown).

Image classification

- learned pos embeddings = no hand-crafted information about the structure
- => pretty cool
- also permutation invariant

| Model | Pretrained? | Top-1 Acc. |
|---|-------------|------------|
| ConvNet baselines | | |
| ResNet-50 [30] | N | 78.6 |
| NFNet-F6+SAM [8] | N | 86.5 |
| Meta Pseudo Labels [60] | Y | 90.2 |
| ViT baselines | | |
| ViT-B/16 [24] | N | 77.9 |
| ViT-H/14 (pretrained) [24] | Y | 88.6 |
| DeiT 1000 epochs [86] | N | 85.2 |
| CaiT-M48 448 [87] | N | 86.5 |
| w/ 2D Fourier features | | |
| Perceiver | N | 78.6 |
| Perceiver IO | N | 79.0 |
| Perceiver IO (pretrained) | Y | 84.5 |
| w/ learned position features | | |
| Perceiver (learned pos) | N | 67.6 |
| Perceiver IO (learned pos) | N | 72.7 |
| w/ 2D conv + maxpool preprocessing | | |
| Perceiver (conv) | N | 77.4 |
| Perceiver IO (conv) | N | 82.1 |

Conclusion

Pros:



- latent space transformer
- linear scaling with input size
- multiple domains, tasks supported

Cons:

- does not support generative modeling

Positional encodings = these are domain dependent

- spatial information is important, Transformer are invariant to it
- => use positional encodings = **Scalable Fourier features**
 - get k frequency bands, k -th band has frequency 2^k
 - sample positional encodings along those bands
 - concat these encodings to the input token vectors
 - 1 row on the figure below = 1 positional encoding



Sources

- Paper: <https://arxiv.org/abs/2107.14795>
- Pytorch: <https://github.com/lucidrains/perceiver-pytorch>
- JAX: <https://github.com/deepmind/deepmind-research/tree/master/perceiver>